

Assignment 2 Report

Github Repo:

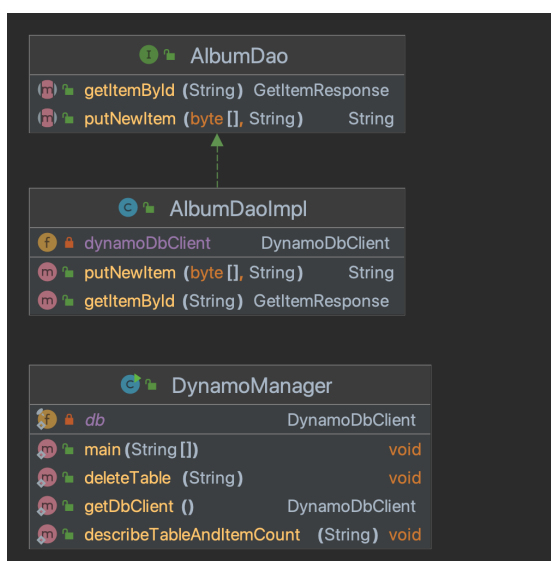
<https://github.com/peihuan-lin/cs6650/tree/main/assignment2>

Description of Data Model

The application employ Amazon DynamoDB as its database for storing album information, including metadata and identifiers for album objects.

- **Album Information Storage:** Each album is uniquely identified by a UUID. Associated with this identifier is album metadata, stored in JSON format, and binary data for the image file.
 - id: the primary key, generated by UUID
 - imageFile: the album image converted to byte, image size is downgrade to 500 byte to keep costs down.
 - profile: the artist, title, and year of the album in JSON format
- **Data Retrieval:** The `getItemById` method enables fetching album data by its unique identifier, facilitating quick lookups.
- **Data Insertion:** The `putNewItem` method allows for inserting new albums, automatically generating a unique identifier and storing the image and profile as binary and string data types, respectively.
- **Database Management:** Encapsulate in `DynamoManager` class to provide utility functions, such as connecting and deleting tables within DynamoDB.

<input type="checkbox"/>	id (String)	imageFile	profile
<input type="checkbox"/>	29ac4bb1-f5f4-4209-...	IVBORwOK...	{"artist": "Sex Pistols", "title": "Never Mind the Bollocks", "year": "1977"}
<input type="checkbox"/>	5a9d5330-880a-4a83-...	IVBORwOK...	{"artist": "Sex Pistols", "title": "Never Mind the Bollocks", "year": "1977"}
<input type="checkbox"/>	81258b82-bd36-44fd-...	IVBORwOK...	{"artist": "Sex Pistols", "title": "Never Mind the Bollocks", "year": "1977"}



Single Server

Configuration

1. **Initial Setup with EC2 t2.micro:** The starting configuration with 50 RCU and WCU led to a high CPU usage of 86% and timeouts during peak loads.
2. **First Adjustment to RCU/WCU:** Upon increasing the provisioned RCU and WCU to 500 each, the system achieved full CPU utilization of 100%, handling 597 requests per second under a load of 10 thread groups with 10 threads each.
3. **Modification for Comparable Testing:** To prevent CPU bottlenecks and obtain a fair comparison of system performance, a reduction in CPU usage was considered essential.
4. **Upgrade to t2.medium Instance:** Retaining the same database settings but moving to a t2.medium instance, the system's throughput improved to 881 requests per second, although the CPU peak utilization remained high at 96%.
5. **Switch to t3.small Instance:** Finally, with the transition to a t3.small instance, there was a balance between throughput and CPU efficiency. The system managed 699 requests per second with a peak CPU utilization of 90%.

Below is the tabulated performance under varying loads with the final configuration on the t3.small instance, maintaining 500 RCU and WCU:

numThreads	threadGroup	numOfRequest	throughput (req/s)	consumed read (units/s)	consumed write (units/s)	CPU utilization (%)
10	10	101000	699	388	403	90
10	20	201000	878	460	921	89
10	30	301000	999	524	1050	93

Output windows for a single server

```
Test load:
threadGroupSize: 10, numThreadGroups: 10, delay: 2
Time taken: 144502 ms
Number of successful requests: 101000
Number of fail requests: 0
Walltime: 144.484 seconds
Total throughput: 699.0393399961241 req/s

Metrics for GET:
Mean response time: 61.82335643564357 ms
Median response time: 52.0 ms
99th response time: 201.00099999999946 ms
Min response time: 13.0 ms
Max response time: 757.0 ms

Metrics for POST:
Mean response time: 80.52737623762376 ms
Median response time: 67.0 ms
99th response time: 260.00099999999946 ms
Min response time: 16.0 ms
Max response time: 3350.0 ms
```

```
Test load:
threadGroupSize: 10, numThreadGroups: 20, delay: 2
Time taken: 228800 ms
Number of successful requests: 201000
Number of fail requests: 0
Walltime: 228.762 seconds
Total throughput: 878.6424318724264 req/s

Metrics for GET:
Mean response time: 98.55069651741293 ms
Median response time: 83.0 ms
99th response time: 339.00100000000094 ms
Min response time: 12.0 ms
Max response time: 994.0 ms

Metrics for POST:
Mean response time: 127.26381592039802 ms
Median response time: 108.0 ms
99th response time: 442.00100000000094 ms
Min response time: 16.0 ms
Max response time: 1252.0 ms
```

```
Test load:
threadGroupSize: 10, numThreadGroups: 30, delay: 2
Time taken: 301264 ms
Number of successful requests: 301000
Number of fail requests: 0
Walltime: 301.206 seconds
Total throughput: 999.3160826809558 req/s

Metrics for GET:
Mean response time: 132.23155149501662 ms
Median response time: 110.0 ms
99th response time: 469.00100000000094 ms
Min response time: 12.0 ms
Max response time: 1349.0 ms

Metrics for POST:
Mean response time: 165.1411096345515 ms
Median response time: 134.0 ms
99th response time: 598.0010000000009 ms
Min response time: 16.0 ms
Max response time: 1367.0 ms
```

Two Servers with Load Balancer

Load Balance & Target Group Configuration

Employing a pair of t3.small EC2 instances grouped within a target group, where a load balancer directs incoming requests between the two.

Load balancer type

Application

Status

Active

VPC

vpc-0a7857a6bc5505472

IP address type

IPv4

Scheme

Internet-facing

Hosted zone

Z1H1FL5HABSF5

Availability Zones

subnet-0fd6cc72b8dfefcb3 us-west-2a (usw2-az2)
subnet-0bf987bd558391d5f us-west-2b (usw2-az1)

Date created

November 2, 2023, 10:38 (UTC-07:00)

Load balancer ARN

arn:aws:elasticloadbalancing:us-west-2:609640301360:loadbalancer/app/alb/4b05a432e622a671

DNS name info

alb-1168581154.us-west-2.elb.amazonaws.com (A Record)

Listeners and rules

Network mapping

Security

Monitoring

Integrations

Attributes

Tags

Listeners and rules (1)

Manage rules

Manage listener

Add listener

A listener checks for connection requests on its configured protocol and port. Traffic received by the listener is routed according to the default action and any additional rules.

Filter listeners

< 1 >

Protocol:Port

HTTP:80

Default action

Forward to target group

- target-group: 1 (100%)
- Group-level stickiness: Off

Rules

1 rule

ARN

ARN

Security policy

Not applicable

Default

Not appl

Details

arn:aws:elasticloadbalancing:us-west-2:609640301360:targetgroup/target-group/a4d4671766320d473

Target type

Instance

Protocol : Port

HTTP: 80

Protocol version

HTTP1

VPC

vpc-0a7857a6bc5505472

IP address type

IPv4

Load balancer

alb

Total targets

2

Healthy

2

Unhealthy

0

Unused

0

Initial

0

Draining

0

Distribution of targets by Availability Zone (AZ)

Select values in this table to see corresponding filters applied to the Registered targets table below.

Targets

Monitoring

Health checks

Attributes

Tags

Registered targets (2)

Deregister

Register target

Filter targets

< 1 >

Instance ID

Name

Port

Zone

Health status

Health status details

I-035a31eda947b0...

server1-t3

8080

us-west-2a

Healthy

I-0b5fe926b6b6686...

server2-t3

8080

us-west-2a

Healthy

numThreads	threadGroup	numOfRequest	throughput (req/s)	consumed read (units/s)	consumed write (units/s)	CPU utilization (%)
10	10	101000	1701	767	1535	61
10	20	201000	1795	996	1932	79
10	30	301000	974	905	1809	78

- For the same load of 10 threads in 10 thread groups, the load balancer increases the throughput to 1701 requests per second with only 61% CPU utilization. This indicates a significant performance enhancement, as the two servers together can handle more than double the requests of a single server at a lower CPU utilization.
- When increasing to 20 thread groups, the throughput further improves to 1795 requests per second, although CPU utilization climbs to 79%, which is still lower than the single server's peak of 93%.
- At 30 thread groups is a reduction in throughput with a CPU utilization of 78%. This decrease could be due to reaching a bottleneck.

Output windows for a two load balanced servers

```
Test load:
threadGroupSize: 10, numThreadGroups: 10, delay: 2
Time taken: 59394 ms
Number of successful requests: 101000
Number of fail requests: 0
Walltime: 59.376 seconds
Total throughput: 1701.0239827539747 req/s

Metrics for GET:
Mean response time: 26.670138613861386 ms
Median response time: 23.0 ms
99th response time: 77.00099999999948 ms
Min response time: 13.0 ms
Max response time: 616.0 ms

Metrics for POST:
Mean response time: 31.777555606149093 ms
Median response time: 27.0 ms
99th response time: 94.07799999999988 ms
Min response time: 16.0 ms
Max response time: 761.0 ms
```

```
Test load:
threadGroupSize: 10, numThreadGroups: 20, delay: 2
Time taken: 111982 ms
Number of successful requests: 201000
Number of fail requests: 0
Walltime: 111.944 seconds
Total throughput: 1795.5406274565853 req/s

Metrics for GET:
Mean response time: 49.458407960199004 ms
Median response time: 40.0 ms
99th response time: 182.00100000000094 ms
Min response time: 13.0 ms
Max response time: 1849.0 ms

Metrics for POST:
Mean response time: 60.07311224540561 ms
Median response time: 47.0 ms
99th response time: 238.09500000000116 ms
Min response time: 16.0 ms
Max response time: 1838.0 ms
```

```
Test load:
threadGroupSize: 10, numThreadGroups: 30, delay: 2
Time taken: 308958 ms
Number of successful requests: 301000
Number of fail requests: 0
Walltime: 308.9 seconds
Total throughput: 974.4253803820008 req/s

Metrics for GET:
Mean response time: 81.77564784053156 ms
Median response time: 49.0 ms
99th response time: 387.01100000001026 ms
Min response time: 13.0 ms
Max response time: 5959.0 ms

Metrics for POST:
Mean response time: 157.49189882077115 ms
Median response time: 62.0 ms
99th response time: 3915.071999999997 ms
Min response time: 16.0 ms
Max response time: 9911.0 ms
```

Set up screenshots

- instance security group

Inbound rules (5)										<div>Manage tags</div> <div>Edit inbound rules</div>	
<div>Filter security group rules</div>										<div>< 1 ></div>	
<input type="checkbox"/>	Name	Security group rul...	IP version	Type	Protocol	Port range	Source	Description			
<input type="checkbox"/>	-	sgr-0a99a556762fa54...	IPv4	HTTPS	TCP	443	0.0.0.0/0	-			
<input type="checkbox"/>	-	sgr-0c5ca86da150536...	IPv4	SSH	TCP	22	23.252.62.110/32	-			
<input type="checkbox"/>	-	sgr-00c01cf9edba5c105	IPv4	SSH	TCP	22	172.92.179.119/32	-			
<input type="checkbox"/>	-	sgr-049868c7ebb723...	-	Custom TCP	TCP	8080	sg-01edf2c89d1a5eb...	ALB			
<input type="checkbox"/>	-	sgr-067d61cf831fea5ea	IPv4	HTTP	TCP	80	0.0.0.0/0	-			

- load balancer security group

Outbound rules (2)										<div>Manage tags</div> <div>Edit outbound rules</div>	
<div>Filter security group rules</div>										<div>< 1 ></div>	
<input type="checkbox"/>	Name	Security group rul...	IP version	Type	Protocol	Port range	Destination	Description			
<input type="checkbox"/>	-	sgr-06bee9d41bde32...	IPv4	HTTP	TCP	80	0.0.0.0/0	health check			
<input type="checkbox"/>	-	sgr-02a1bc071967e1...	-	Custom TCP	TCP	8080	sg-015fb8551dffe6e1...	instance			

Tune the System

Adjusted based on a load of 10 threads in 30 thread groups.

DataBase Improvements

The transition to on-demand capacity yielded a marked improvement in system throughput, increased from 974 to 1922 requests per second. The adoption of on-demand capacity enabled the system to better harness DynamoDB's automatic scaling capabilities, ensuring robust performance and resource optimization in response to high loads.

Configuration	Throughput (req/s)	Consumed Read (units/s)	Consumed Write (units/s)	CPU Utilization (%)
provisioned	974	905	1809	78
on demand	1922	973	1948	81

```
Test load:
threadGroupSize: 10, numThreadGroups: 30, delay: 2
Time taken: 156659 ms
Number of successful requests: 301000
Number of fail requests: 0
Walltime: 156.601 seconds
Total throughput: 1922.08223446849 req/s

Metrics for GET:
Mean response time: 68.0457607973422 ms
Median response time: 54.0 ms
99th response time: 270.00100000000094 ms
Min response time: 13.0 ms
Max response time: 901.0 ms

Metrics for POST:
Mean response time: 85.1850845340412 ms
Median response time: 66.0 ms
99th response time: 350.09799999999814 ms
Min response time: 17.0 ms
Max response time: 1001.0 ms
```

Set up screenshots

Edit read/write capacity

Capacity mode [Info](#)

☒ On-demand
Simplify billing by paying for the actual reads and writes your application performs.

☐ Provisioned
Manage and optimize your costs by allocating read/write capacity in advance.

Server Improvements

Increase Instance number

The adoption of on-demand capacity was sustained as I increased the number of t3.small servers.

Configuration	Throughput (req/s)	Consumed Read (units/s)	Consumed Write (units/s)	CPU Utilization (%)
2 servers	1922	973	1948	81
5 servers	3603	1333	2663	56
9 servers	4892	1805	3609	45

- **With 5 t3.small Servers:** The system's throughput escalated to 3603 requests per second, with a marked decrease in CPU utilization to 56%. This indicates that spreading the load across more servers leads to more efficient request handling and less CPU strain.

```
Test load:
threadGroupSize: 10, numThreadGroups: 30, delay: 2
Time taken: 83585 ms
Number of successful requests: 301000
Number of fail requests: 0
Walltime: 83.527 seconds
Total throughput: 3603.6251750930837 req/s

Metrics for GET:
Mean response time: 36.633960132890365 ms
Median response time: 28.0 ms
99th response time: 149.00100000000094 ms
Min response time: 13.0 ms
Max response time: 666.0 ms

Metrics for POST:
Mean response time: 44.945029247536134 ms
Median response time: 33.0 ms
99th response time: 180.05 ms
Min response time: 16.0 ms
Max response time: 762.0 ms
```

- **With 9 t3.small Servers:** The benefits of scaling out became even more evident. Throughput reached 4892 requests per second, CPU utilization further decreased to 45%, reinforcing the effectiveness of this scaling strategy.

```
Test load:
threadGroupSize: 10, numThreadGroups: 30, delay: 2
Time taken: 61583 ms
Number of successful requests: 301000
Number of fail requests: 0
Walltime: 61.525 seconds
Total throughput: 4892.320195042666 req/s

Metrics for GET:
Mean response time: 27.556644518272424 ms
Median response time: 23.0 ms
99th response time: 104.00100000000093 ms
Min response time: 13.0 ms
Max response time: 617.0 ms

Metrics for POST:
Mean response time: 32.65234802688768 ms
Median response time: 27.0 ms
99th response time: 123.09899999999907 ms
Min response time: 16.0 ms
Max response time: 649.0 ms
```

Set up screenshots

Target groups (3) Info

Filter target groups

<input type="checkbox"/>	Name	ARN	Port	Protocol	Target type	Load balancer	VPC ID
<input type="checkbox"/>	fiveT3Server	arn:aws:elasti...	80	HTTP	Instance	None associated	vpc-0a7857a6bc5505472
<input type="checkbox"/>	nineT3Server	arn:aws:elasti...	80	HTTP	Instance	None associated	vpc-0a7857a6bc5505472
<input type="checkbox"/>	twoT3Server	arn:aws:elasti...	80	HTTP	Instance	None associated	vpc-0a7857a6bc5505472

Load balancers (3)

Elastic Load Balancing scales your load balancer capacity automatically in response to changes in incoming traffic.

Filter load balancers

<input type="checkbox"/>	Name	DNS name	State	VPC ID	Availability Zones	
<input type="checkbox"/>	alb-2	alb-2-1194855567.us-wes...	Active	vpc-0a7857a6bc5505...	2 Availability Zones	a
<input type="checkbox"/>	alb-5	alb-5-1727125722.us-wes...	Active	vpc-0a7857a6bc5505...	2 Availability Zones	a
<input type="checkbox"/>	alb-9	alb-9-912815484.us-west-...	Active	vpc-0a7857a6bc5505...	2 Availability Zones	a

alb-9

Details

Load balancer type

Application

Status

Active

VPC

[vpc-0a7857a6bc5505472](#)

IP address type

IPv4

Scheme

Internet-facing

Hosted zone

Z1H1FL5HABSF5

Availability Zones

[subnet-0fd6cc72b8dfefcb3](#) us-west-2a (usw2-az2)
[subnet-0bf987bd558391d5f](#) us-west-2b (usw2-az1)

Date created

November 4, 2023, 21:16 (UTC-07:00)

Load balancer ARN

[arn:aws:elasticloadbalancing:us-west-2:609640301360:loadbalancer/app/alb-9/55db697db36e4b1d](#)

DNS name

[alb-9-912815484.us-west-2.elb.amazonaws.com](#) (A Record)

Listeners and rules

Network mappingSecurityMonitoringIntegrationsAttributesTags

Listeners and rules (1) Info

Filter listeners

<input type="checkbox"/>	Protocol:Port	Default action	Rules	ARN	Security policy	Default SSL/TLS certificate	Tags
<input type="checkbox"/>	HTTP:80	<div>Forward to target group<ul style="list-style-type: none">nineT3Server: 1 (100%)Group-level stickiness: Off</div>	1 rule	ARN	Not applicable	Not applicable	0 tags

Overall Throughput Improvement

- Moving from a single server to two load-balanced servers, then optimizing to on-demand capacity and scaling to 9 servers, the system's throughput improved significantly:
 - From 999 req/s (single server) to 4892 req/s (9 load balanced servers), resulting in a **390% improvement** in throughput.
 - CPU utilization was optimized from a peak of 93% (single server) to 45% (9 servers), showing a more efficient system.

DB Config	Server Config	Throughput (req/s)	Consumed Read (units/s)	Consumed Write (units/s)	CPU Utilization (%)
Provisioned	single server	999	524	1050	93
Provisioned	2 load balanced server	974	905	1809	78
On-demand	2 load balanced server	1922	973	1948	81
On-demand	5 load balanced server	3603	1333	2663	56
On-demand	9 load balanced server	4892	1805	3609	45